

5. Observational Studies

ISS5096 || ECI

Jaewon (“Jay-one”) Yoo

National Tsing Hua University

Roadmap

1. Identification in Observational Studies
2. Selection on Observables
3. Sensitivity Analysis
4. Partial Identification

Where are we? Where are we going?

- So far: experiments where design makes things easier.

Where are we? Where are we going?

- So far: experiments where design makes things easier.
- Today: what happens when we have observational studies to work with?
 - **Identification** and **selection on observables**: assumptions and estimation.

Where are we? Where are we going?

- So far: experiments where design makes things easier.
- Today: what happens when we have observational studies to work with?
 - **Identification** and **selection on observables**: assumptions and estimation.
 - What if the assumptions are wrong? $\rightsquigarrow$ **sensitivity analysis** and **partial identification**.
 - Rest of the course will cover different designs for observational studies.

Where are we? Where are we going?

- So far: experiments where design makes things easier.
- Today: what happens when we have observational studies to work with?
 - **Identification** and **selection on observables**: assumptions and estimation.
 - What if the assumptions are wrong? $\rightsquigarrow$ **sensitivity analysis** and **partial identification**.
 - Rest of the course will cover different designs for observational studies.
- Q: Why are observational studies in causal inference important? (What are the limitations of RCTs?)

Where are we? Where are we going?



Source: Twitter @NobelPrize

Where are we? Where are we going?



Scientific Background on the Sveriges Riksbank Prize in Economic Sciences in Memory of Alfred Nobel 2021

ANSWERING CAUSAL QUESTIONS USING OBSERVATIONAL DATA

The Committee for the Prize in Economic Sciences in Memory of Alfred Nobel

*“Taken together, therefore, the Laureates’ contributions have played a central role in establishing the so-called design-based approach in economics. This approach – aimed at **emulating a randomized experiment to answer a causal question using observational data** – has transformed applied work and improved researchers’ ability to answer causal questions of great importance for economic and social policy using observational data.” (p.2)*

1/ Identification in Observational Studies

Randomized Experiment Review

- **Experiment:** when the researcher controls the treatment assignment.
 - $p_i = \mathbb{P}[D_i = 1]$ is the probability of treatment assignment.
 - p_i is controlled by & known to researcher in an experiment.

Randomized Experiment Review

- **Experiment:** when the researcher controls the treatment assignment.
 - $p_i = \mathbb{P}[D_i = 1]$ is the probability of treatment assignment.
 - p_i is controlled by & known to researcher in an experiment.
- **Randomized experiment** is an experiment with two properties:
 1. **Positivity:** assignment is probabilistic (and not deterministic):
 $0 < \mathbb{P}[D_i = 1] < 1$

Randomized Experiment Review

- **Experiment:** when the researcher controls the treatment assignment.
 - $p_i = \mathbb{P}[D_i = 1]$ is the probability of treatment assignment.
 - p_i is controlled by & known to researcher in an experiment.
- **Randomized experiment** is an experiment with two properties:
 1. **Positivity:** assignment is probabilistic (and not deterministic):
 $0 < \mathbb{P}[D_i = 1] < 1$
 2. **Unconfoundedness:** $\mathbb{P}[D_i = 1 | \mathbf{Y}(1), \mathbf{Y}(0)] = \mathbb{P}[D_i = 1]$
 - Treatment assignment does not depend on any potential outcomes.
 - Sometimes written as $D_i \perp\!\!\!\perp (\mathbf{Y}(1), \mathbf{Y}(0))$.

What is the Selection Problem?

- What if we **observe** a non-randomized treatment?
 - Maybe treatment assignment is **confounded** so D_i is related to POs.

What is the Selection Problem?

- What if we **observe** a non-randomized treatment?
 - Maybe treatment assignment is **confounded** so D_i is related to POs.
- What can we learn about the ATE here? Look at the diff-in-means:

$$\mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0]$$

What is the Selection Problem?

- What if we **observe** a non-randomized treatment?
 - Maybe treatment assignment is **confounded** so D_i is related to POs.
- What can we learn about the ATE here? Look at the diff-in-means:

$$\begin{aligned} & \mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0] \\ &= \mathbb{E}[Y_i(1)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0] \quad (\text{consistency}) \end{aligned}$$

What is the Selection Problem?

- What if we **observe** a non-randomized treatment?
 - Maybe treatment assignment is **confounded** so D_i is related to POs.
- What can we learn about the ATE here? Look at the diff-in-means:

$$\begin{aligned} & \mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0] \\ &= \mathbb{E}[Y_i(1)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0] \\ &= \mathbb{E}[Y_i(1)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 1] + \mathbb{E}[Y_i(0)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0] \end{aligned}$$

What is the Selection Problem?

- What if we **observe** a non-randomized treatment?
 - Maybe treatment assignment is **confounded** so D_i is related to POs.
- What can we learn about the ATE here? Look at the diff-in-means:

$$\begin{aligned} & \mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0] \\ &= \mathbb{E}[Y_i(1)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0] \\ &= \mathbb{E}[Y_i(1)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 1] + \mathbb{E}[Y_i(0)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0] \\ &= \underbrace{\mathbb{E}[Y_i(1) - Y_i(0)|D_i = 1]}_{\text{ATT}} + \underbrace{\mathbb{E}[Y_i(0)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0]}_{\text{selection bias}} \end{aligned}$$

What is the Selection Problem?

- What if we **observe** a non-randomized treatment?
 - Maybe treatment assignment is **confounded** so D_i is related to POs.
- What can we learn about the ATE here? Look at the diff-in-means:

$$\begin{aligned} & \mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0] \\ &= \mathbb{E}[Y_i(1)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0] \\ &= \mathbb{E}[Y_i(1)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 1] + \mathbb{E}[Y_i(0)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0] \\ &= \underbrace{\mathbb{E}[Y_i(1) - Y_i(0)|D_i = 1]}_{\text{ATT}} + \underbrace{\mathbb{E}[Y_i(0)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0]}_{\text{selection bias}} \end{aligned}$$

- Without unconfoundedness: naive diff-in-means = PATT + selection bias
- **Selection bias:** how different the treated and control groups are in terms of their potential outcome under control.

Selection Bias = Unidentified ATT

$$\mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0] = \underbrace{\tau_{\text{treated}}}_{\text{ATT}} + \underbrace{\mathbb{E}[Y_i(0)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0]}_{\text{selection bias}}$$

- Difference in means: a combination of two unknown quantities.
 - Can't distinguish if a diff-in-means is the ATT or selection bias.

Selection Bias = Unidentified ATT

$$\mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0] = \underbrace{\tau_{\text{treated}}}_{\text{ATT}} + \underbrace{\mathbb{E}[Y_i(0)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0]}_{\text{selection bias}}$$

- Difference in means: a combination of two unknown quantities.
 - Can't distinguish if a diff-in-means is the ATT or selection bias.
- Example: effect of comment sections on support for online influencers.
 - Naive estimate: influencers do worse without comment sections than with them.
 - $\rightsquigarrow$ negative ATT **OR** positive ATT with large negative selection bias.
 - SB = influencers that disable user comments are worse than those that keep them, even if they posted the same content.

Selection Bias = Unidentified ATT

$$\mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0] = \underbrace{\tau_{\text{treated}}}_{\text{ATT}} + \underbrace{\mathbb{E}[Y_i(0)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0]}_{\text{selection bias}}$$

- Difference in means: a combination of two unknown quantities.
 - Can't distinguish if a diff-in-means is the ATT or selection bias.
- Example: effect of comment sections on support for online influencers.
 - Naive estimate: influencers do worse without comment sections than with them.
 - $\rightsquigarrow$ negative ATT **OR** positive ATT with large negative selection bias.
 - SB = influencers that disable user comments are worse than those that keep them, even if they posted the same content.
- With an unbounded Y_i , we cannot even bound the ATT because, in principle, SB could be anywhere from $-\infty$ to ∞ .

Selection Bias = Unidentified ATT

$$\mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0] = \underbrace{\tau_{\text{treated}}}_{\text{ATT}} + \underbrace{\mathbb{E}[Y_i(0)|D_i = 1] - \mathbb{E}[Y_i(0)|D_i = 0]}_{\text{selection bias}}$$

- Difference in means: a combination of two unknown quantities.
 - Can't distinguish if a diff-in-means is the ATT or selection bias.
- Example: effect of comment sections on support for online influencers.
 - Naive estimate: influencers do worse without comment sections than with them.
 - $\rightsquigarrow$ negative ATT **OR** positive ATT with large negative selection bias.
 - SB = influencers that disable user comments are worse than those that keep them, even if they posted the same content.
- With an unbounded Y_i , we cannot even bound the ATT because, in principle, SB could be anywhere from $-\infty$ to ∞ .
- We say ATT (as well as ATE) are **unidentified** w/o further assumptions.

What is identification?

- **Identification** connects the counterfactual to the observed.
 - **Counterfactual distribution** $\mathbb{P}^*$ of $\{Y_i(1), Y_i(0), D_i, \mathbf{X}_i\}$.
 - **Observational distribution** $\mathbb{P}$ of $\{Y_i, D_i, \mathbf{X}_i\}$.
 - Causal quantities are functions of $\mathbb{P}^*$, but we get samples from $\mathbb{P}$.
 - $\rightsquigarrow$ We can only learn about $\mathbb{P}^*$ through $\mathbb{P}$!

What is identification?

- **Identification** connects the counterfactual to the observed.
 - **Counterfactual distribution** $\mathbb{P}^*$ of $\{Y_i(1), Y_i(0), D_i, \mathbf{X}_i\}$.
 - **Observational distribution** $\mathbb{P}$ of $\{Y_i, D_i, \mathbf{X}_i\}$.
 - Causal quantities are functions of $\mathbb{P}^*$, but we get samples from $\mathbb{P}$.
 - $\rightsquigarrow$ We can only learn about $\mathbb{P}^*$ through $\mathbb{P}$!
- Quantity ψ ($/(\text{p})\text{sal}/$) is **identified** if we can write it as function of $\mathbb{P}$.
 - Would we know this quantity if we had access to unlimited data?
 - $\rightsquigarrow$ no worrying about estimation uncertainty here.

What is identification?

- **Identification** connects the counterfactual to the observed.
 - **Counterfactual distribution** $\mathbb{P}^*$ of $\{Y_i(1), Y_i(0), D_i, \mathbf{X}_i\}$.
 - **Observational distribution** $\mathbb{P}$ of $\{Y_i, D_i, \mathbf{X}_i\}$.
 - Causal quantities are functions of $\mathbb{P}^*$, but we get samples from $\mathbb{P}$.
 - $\rightsquigarrow$ We can only learn about $\mathbb{P}^*$ through $\mathbb{P}$!
- Quantity ψ ($\psi(p)$) is **identified** if we can write it as function of $\mathbb{P}$.
 - Would we know this quantity if we had access to unlimited data?
 - $\rightsquigarrow$ no worrying about estimation uncertainty here.
- Connecting counterfactuals to the observational requires **assumptions**.
 - “**What is your identification strategy?**” = what are the assumptions that allow you to claim that you’ve estimated a causal effect?
 - Research design can help justify assumptions (experiments, RDD, etc).
 - Or you will need to justify them through argument.

Identification vs. Estimation

- Identification tells us **what** to estimate, not **how**.
 - If identified, we know our causal parameter is some function of $\mathbb{P}$.

Identification vs. Estimation

- Identification tells us **what** to estimate, not **how**.
 - If identified, we know our causal parameter is some function of $\mathbb{P}$.
 - For example, let's consider the **population** diff-in-means:

$$\mathbb{E}[Y_i | D_i = 1] - \mathbb{E}[Y_i | D_i = 0]$$

- But, $\mathbb{P}$ is not directly observable since it's a population distribution!

Identification vs. Estimation

- Identification tells us **what** to estimate, not **how**.
 - If identified, we know our causal parameter is some function of $\mathbb{P}$.
 - For example, let's consider the **population** diff-in-means:

$$\mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0]$$

- But, $\mathbb{P}$ is not directly observable since it's a population distribution!
- Once identified, we need to actually **estimate** the function of $\mathbb{P}$.
 - $\hat{\tau}_{\text{diff}}$ is an estimator for population diff-in-means
 - Now just estimating conditional expectations, etc.
 - $\rightsquigarrow$ **after identification, causal inference part done**
 - Purely a statistical question from here on out.

Identification vs. Estimation

- Identification tells us **what** to estimate, not **how**.
 - If identified, we know our causal parameter is some function of $\mathbb{P}$.
 - For example, let's consider the **population** diff-in-means:

$$\mathbb{E}[Y_i|D_i = 1] - \mathbb{E}[Y_i|D_i = 0]$$

- But, $\mathbb{P}$ is not directly observable since it's a population distribution!
- Once identified, we need to actually **estimate** the function of $\mathbb{P}$.
 - $\hat{\tau}_{\text{diff}}$ is an estimator for population diff-in-means
 - Now just estimating conditional expectations, etc.
 - $\rightsquigarrow$ **after identification, causal inference part done**
 - Purely a statistical question from here on out.
- **Identification** comes first, then comes **estimation**.
 - Without identification, properties of the estimator are unimportant.
 - keep them separate: estimator shouldn't drive identification.

What is Confounding?

- **Confounding:** treatment and potential outcomes are not independent!
 - Due to “common causes” of Y_i and D_i .
 - Main concern in observational studies.

What is Confounding?

- **Confounding:** treatment and potential outcomes are not independent!
 - Due to “common causes” of Y_i and D_i .
 - Main concern in observational studies.
- Pervasive in management/social sciences:
 - Effect of job training program on employment (confounder: motivation)

What is Confounding?

- **Confounding:** treatment and potential outcomes are not independent!
 - Due to “common causes” of Y_i and D_i .
 - Main concern in observational studies.
- Pervasive in management/social sciences:
 - Effect of job training program on employment (confounder: motivation)
 - Effect of college GPAs on salary (confounder: intelligence)

What is Confounding?

- **Confounding:** treatment and potential outcomes are not independent!
 - Due to “common causes” of Y_i and D_i .
 - Main concern in observational studies.
- Pervasive in management/social sciences:
 - Effect of job training program on employment (confounder: motivation)
 - Effect of college GPAs on salary (confounder: intelligence)
 - Effect of income on voting (confounder: age)

What is Confounding?

- **Confounding:** treatment and potential outcomes are not independent!
 - Due to “common causes” of Y_i and D_i .
 - Main concern in observational studies.
- Pervasive in management/social sciences:
 - Effect of job training program on employment (confounder: motivation)
 - Effect of college GPAs on salary (confounder: intelligence)
 - Effect of income on voting (confounder: age)
 - Effect of corporate giants on economic development (confounder: previous economic development)
- Confounding $\rightsquigarrow$ incomplete identification of ATE $\rightsquigarrow$ biased estimators.
- What to do?

2/ Selection on Observables

Observational Studies

- Many different types of identification assumptions we'll cover.

Observational Studies

- Many different types of identification assumptions we'll cover.
- Begin with most common observational assumptions:
 1. **No unmeasured confounding:** $\{Y_i(1), Y_i(0)\} \perp\!\!\!\perp D_i | \mathbf{X}_i$
 - Also called: conditional unconfoundedness, weak ignorability, selection on observables, no omitted variables, exogenous, conditional exchangeability, etc.
 - $\rightsquigarrow$ Conditional on some covariates, D_i is (effectively) randomly assigned.
 2. **Positivity or Overlap:** $0 < \mathbb{P}[D_i = 1 | \mathbf{X}_i] < 1$
 - Treatment and control are both possible at every value of $\mathbf{X}_i$
 - $\rightsquigarrow$ There are both treated and untreated units for each level of $\mathbf{X}_i$ (i.e., “common support”).

Observational Studies

- Many different types of identification assumptions we'll cover.
- Begin with most common observational assumptions:
 1. **No unmeasured confounding:** $\{Y_i(1), Y_i(0)\} \perp\!\!\!\perp D_i | \mathbf{X}_i$
 - Also called: conditional unconfoundedness, weak ignorability, selection on observables, no omitted variables, exogenous, conditional exchangeability, etc.
 - $\rightsquigarrow$ Conditional on some covariates, D_i is (effectively) randomly assigned.
 2. **Positivity or Overlap:** $0 < \mathbb{P}[D_i = 1 | \mathbf{X}_i] < 1$
 - Treatment and control are both possible at every value of $\mathbf{X}_i$
 - $\rightsquigarrow$ There are both treated and untreated units for each level of $\mathbf{X}_i$ (i.e., “common support”).
- We'll take $\mathbf{X}_i$ as a ‘given’ for now and see later how we might choose it.

Observational Studies

- Many different types of identification assumptions we'll cover.
- Begin with most common observational assumptions:
 1. **No unmeasured confounding:** $\{Y_i(1), Y_i(0)\} \perp\!\!\!\perp D_i | \mathbf{X}_i$
 - Also called: conditional unconfoundedness, weak ignorability, selection on observables, no omitted variables, exogenous, conditional exchangeability, etc.
 - $\rightsquigarrow$ Conditional on some covariates, D_i is (effectively) randomly assigned.
 2. **Positivity** or **Overlap:** $0 < \mathbb{P}[D_i = 1 | \mathbf{X}_i] < 1$
 - Treatment and control are both possible at every value of $\mathbf{X}_i$
 - $\rightsquigarrow$ There are both treated and untreated units for each level of $\mathbf{X}_i$ (i.e., “common support”).
- We'll take $\mathbf{X}_i$ as a ‘given’ for now and see later how we might choose it.
- These are assumptions that **can be wrong!**

Identification of the ATE

- Positivity and no unmeasured confounders will identify the PATE:

$$\tau = \mathbb{E}[Y_i(1) - Y_i(0)]$$

Identification of the ATE

- Positivity and no unmeasured confounders will identify the PATE:

$$\begin{aligned}\tau &= \mathbb{E}[Y_i(1) - Y_i(0)] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1) - Y_i(0)|\mathbf{x}_i]]\end{aligned}$$

Identification of the ATE

- Positivity and no unmeasured confounders will identify the PATE:

$$\begin{aligned}\tau &= \mathbb{E}[Y_i(1) - Y_i(0)] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1) - Y_i(0)|\mathbf{x}_i]] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1)|\mathbf{x}_i] - \mathbb{E}[Y_i(0)|\mathbf{x}_i]]\end{aligned}$$

Identification of the ATE

- Positivity and no unmeasured confounders will identify the PATE:

$$\begin{aligned}\tau &= \mathbb{E}[Y_i(1) - Y_i(0)] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1) - Y_i(0)|\mathbf{x}_i]] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1)|\mathbf{x}_i] - \mathbb{E}[Y_i(0)|\mathbf{x}_i]] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1)|D_i = 1, \mathbf{x}_i] - \mathbb{E}[Y_i(0)|D_i = 0, \mathbf{x}_i]]\end{aligned}$$

Identification of the ATE

- Positivity and no unmeasured confounders will identify the PATE:

$$\begin{aligned}\tau &= \mathbb{E}[Y_i(1) - Y_i(0)] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1) - Y_i(0)|\mathbf{x}_i]] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1)|\mathbf{x}_i] - \mathbb{E}[Y_i(0)|\mathbf{x}_i]] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1)|D_i = 1, \mathbf{x}_i] - \mathbb{E}[Y_i(0)|D_i = 0, \mathbf{x}_i]] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i|D_i = 1, \mathbf{x}_i] - \mathbb{E}[Y_i|D_i = 0, \mathbf{x}_i]]\end{aligned}$$

- Useful to write the treated and control CEFs:

$$\mu_1(\mathbf{x}) = \mathbb{E}[Y_i(1)|\mathbf{x}_i = \mathbf{x}], \quad \mu_0(\mathbf{x}) = \mathbb{E}[Y_i(0)|\mathbf{x}_i = \mathbf{x}]$$

- How the mean of the potential outcomes vary with the covariates.

Identification of the ATE

- Positivity and no unmeasured confounders will identify the PATE:

$$\begin{aligned}\tau &= \mathbb{E}[Y_i(1) - Y_i(0)] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1) - Y_i(0)|\mathbf{x}_i]] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1)|\mathbf{x}_i] - \mathbb{E}[Y_i(0)|\mathbf{x}_i]] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i(1)|D_i = 1, \mathbf{x}_i] - \mathbb{E}[Y_i(0)|D_i = 0, \mathbf{x}_i]] \\ &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}[Y_i|D_i = 1, \mathbf{x}_i] - \mathbb{E}[Y_i|D_i = 0, \mathbf{x}_i]]\end{aligned}$$

- Useful to write the treated and control CEFs:

$$\mu_1(\mathbf{x}) = \mathbb{E}[Y_i(1)|\mathbf{x}_i = \mathbf{x}], \quad \mu_0(\mathbf{x}) = \mathbb{E}[Y_i(0)|\mathbf{x}_i = \mathbf{x}]$$

- How the mean of the potential outcomes vary with the covariates.
- Key part of the identification above:

$$\underbrace{\mu_1(\mathbf{x})}_{\text{counterfactual}} = \underbrace{\mathbb{E}[\textcolor{red}{Y}_i|D_i = 1, \mathbf{x}_i = \mathbf{x}]}_{\text{observational}}, \quad \mu_0(\mathbf{x}) = \mathbb{E}[\textcolor{red}{Y}_i|D_i = 0, \mathbf{x}_i = \mathbf{x}]$$

Regression Estimation of the ATE

- Identification done, now turn to estimation!

Regression Estimation of the ATE

- Identification done, now turn to estimation!
- Regression estimators $\widehat{\mu}_1(\mathbf{x})$ and $\widehat{\mu}_0(\mathbf{x})$.
 - Might be linear or nonlinear models.
 - Safest practice: estimate separate regressions in each treatment group.
 - Sometimes called an **imputation** or **plug-in** estimator.

Regression Estimation of the ATE

- Identification done, now turn to estimation!
- Regression estimators $\widehat{\mu}_1(\mathbf{x})$ and $\widehat{\mu}_0(\mathbf{x})$.
 - Might be linear or nonlinear models.
 - Safest practice: estimate separate regressions in each treatment group.
 - Sometimes called an **imputation** or **plug-in** estimator.
- Regression/imputation estimator of the ATE:

$$\widehat{\tau}_{\text{reg}} = \frac{1}{n} \sum_{i=1}^n \widehat{\mu}_1(\mathbf{x}_i) - \widehat{\mu}_0(\mathbf{x}_i)$$

Regression Estimation of the ATE

- Identification done, now turn to estimation!
- Regression estimators $\widehat{\mu}_1(\mathbf{x})$ and $\widehat{\mu}_0(\mathbf{x})$.
 - Might be linear or nonlinear models.
 - Safest practice: estimate separate regressions in each treatment group.
 - Sometimes called an **imputation** or **plug-in** estimator.
- Regression/imputation estimator of the ATE:

$$\widehat{\tau}_{\text{reg}} = \frac{1}{n} \sum_{i=1}^n \widehat{\mu}_1(\mathbf{x}_i) - \widehat{\mu}_0(\mathbf{x}_i)$$

- Procedure:
 1. Obtain predicted values for all units when $D_i = 1$.
 2. Obtain predicted values for all units when $D_i = 0$.
 3. Take the average difference between these predicted values.

Coefficients?

$$\widehat{\tau}_{\text{reg}} = \frac{1}{n} \sum_{i=1}^n \widehat{\mu}_1(\mathbf{x}_i) - \widehat{\mu}_0(\mathbf{x}_i)$$

- Under linear models, $\widehat{\tau}_{\text{reg}}$ is sometimes equivalent to a coefficient.

Coefficients?

$$\widehat{\tau}_{\text{reg}} = \frac{1}{n} \sum_{i=1}^n \widehat{\mu}_1(\mathbf{x}_i) - \widehat{\mu}_0(\mathbf{x}_i)$$

- Under linear models, $\widehat{\tau}_{\text{reg}}$ is sometimes equivalent to a coefficient.
- Uninteracted OLS:
 - $\widehat{\mu}_1(x)$ and $\widehat{\mu}_0(x)$ are from the same OLS model $Y \sim D + X$.
 - $\widehat{\tau}_{\text{reg}} \equiv$ estimated coefficient on D_i .

Coefficients?

$$\widehat{\tau}_{\text{reg}} = \frac{1}{n} \sum_{i=1}^n \widehat{\mu}_1(\mathbf{x}_i) - \widehat{\mu}_0(\mathbf{x}_i)$$

- Under linear models, $\widehat{\tau}_{\text{reg}}$ is sometimes equivalent to a coefficient.
- Uninteracted OLS:
 - $\widehat{\mu}_1(x)$ and $\widehat{\mu}_0(x)$ are from the same OLS model $Y \sim D + X$.
 - $\widehat{\tau}_{\text{reg}} \equiv$ estimated coefficient on D_i .
- Fully interacted OLS:
 - $\widehat{\mu}_1(x)$ and $\widehat{\mu}_0(x)$ are from fully interacted OLS with centered covariates.
 - The two CEFs are estimated separately with different slopes for treated and control.
 - $\widehat{\tau}_{\text{reg}} \equiv$ estimated coefficient on D_i .

Coefficients?

$$\widehat{\tau}_{\text{reg}} = \frac{1}{n} \sum_{i=1}^n \widehat{\mu}_1(\mathbf{x}_i) - \widehat{\mu}_0(\mathbf{x}_i)$$

- Under linear models, $\widehat{\tau}_{\text{reg}}$ is sometimes equivalent to a coefficient.
- Uninteracted OLS:
 - $\widehat{\mu}_1(x)$ and $\widehat{\mu}_0(x)$ are from the same OLS model $Y \sim D + X$.
 - $\widehat{\tau}_{\text{reg}} \equiv$ estimated coefficient on D_i .
- Fully interacted OLS:
 - $\widehat{\mu}_1(x)$ and $\widehat{\mu}_0(x)$ are from fully interacted OLS with centered covariates.
 - The two CEFs are estimated separately with different slopes for treated and control.
 - $\widehat{\tau}_{\text{reg}} \equiv$ estimated coefficient on D_i .
- These make two very different assumptions about the CEFs!

Variance Estimation

- How do we get estimates of the variance of $\widehat{\tau}_{\text{reg}}$?

Variance Estimation


- How do we get estimates of the variance of $\widehat{\tau}_{\text{reg}}$?
- If an OLS coefficient $\rightsquigarrow$ use EHW variance estimator.
- Analytic expressions can be derived, but complicated!

Variance Estimation

- How do we get estimates of the variance of $\widehat{\tau}_{\text{reg}}$?
- If an OLS coefficient $\rightsquigarrow$ use EHW variance estimator.
- Analytic expressions can be derived, but complicated!
- Computational alternative: **(nonparametric) bootstrap**
 - Randomly resample n rows of the data with replacement.
 - Refit the regressions on the bootstrapped data.
 - Calculate $\widehat{\tau}_{\text{reg}}$ in each bootstrap.
 - Repeat several times and use empirical variance of the bootstraps.


Imputation Estimator Visualization

```
1 > toy_data <- read_csv("https://bit.ly/3v0y2Ao")
```




Imputation Estimator Visualization

```
1 > lm0 <- lm(y~x, data = toy_data, subset = d==0); lm1 <- lm(y~x, data = toy_data, subset = d==1)
```



Imputation Estimator Visualization

```
1 > mu0.imps = predict(lm0, toy_data); mu1.imps = predict(lm1, toy_data)
2 > cat("Estimate of ATE:", mean(mu1.imps - mu0.imps))
3 Estimate of ATE: 1.285176
```



Fully Interacted OLS & Imputation Estimator

- What if $\hat{\mu}_1(\mathbf{x})$ and $\hat{\mu}_0(\mathbf{x})$ are from fully interacted OLS with centered covariates?
 - Equivalent to running separate models for $\hat{\mu}_1(\mathbf{x})$ and $\hat{\mu}_0(\mathbf{x})$ (i.e., imputation estimator)

```
1 > toy_data$x_tilde <- toy_data$x - mean(toy_data$x)
2 > mod_full <- lm(y ~ d + x_tilde + d * x_tilde, data = toy_data)
3
4 > cat("\nEstimate of ATE (Imputation):", mean(mu1.imps - mu0.imps),
5       "\nEstimated coefficient on Di from full int.", mod_full$coefficients["d"])
6
7 Estimate of ATE (Imputation): 1.285176
8 Estimated coefficient on Di from full int. 1.285176
```

- $\rightsquigarrow$ Recall: Under linear models, $\hat{\tau}_{\text{reg}}$ is sometimes equivalent to a coefficient.
 - $\hat{\tau}_{\text{reg}} \equiv$ estimated coefficient on D_i .
 - Would be the same for uninteracted model, except the variance will be larger (less precision).

Variance Estimation w/ Bootstrap


- Again, how do we get estimates of the variance of $\hat{\tau}_{\text{reg}}$?
 - (Nonparametric) bootstrap: recall source of variance is due to **sampling**
 - Idea: view sample (data) as “population” $\rightsquigarrow$ in-sample “sampling”

```
1 > set.seed(02138); sims<-500; tau_hat_draws<-rep(NA, sims)
2 > for (i in 1:sims) { # Repeat the following several times
3   # 1. Randomly resample n rows of the data with replacement
4   sample_boot <- dplyr::slice_sample(toy_data, n = nrow(toy_data), replace = TRUE)
5
6   # 2. Refit the regressions on the bootstrapped data
7   model <- lm(y ~ d + x_tilde + d*x_tilde, data = toy_data)
8   dat1 <- sample_boot; dat1$d <- 1
9   dat0 <- sample_boot; dat0$d <- 0
10  mu1_hat <- predict(model, newdata = dat1)
11  mu0_hat <- predict(model, newdata = dat0)
12
13  # 3. Calculate tau_hat in each bootstrap
14  tau_hat_draws[i] <- mean(mu1_hat - mu0_hat)
15 }
16
17 > # 4. Use empirical variance of the bootstraps
18 > var(tau_hat_draws)
19 [1] 0.000254049
```

Nonlinear Relationships


- Same idea but with nonlinear relationship between Y_i and X_i :

```
1 > toy_data_02 <- read_csv("https://bit.ly/4c0g0Pl") # sim data with nonlinear relationship
```




Nonlinear Relationships

- Same idea but with nonlinear relationship between Y_i and X_i :



Nonlinear Relationships

- Same idea but with nonlinear relationship between Y_i and X_i :




Using Semiparametric Regression


- Here, CEFs are nonlinear, but we don't know their form.
- We can use generalized additive models (GAMs) from the mgcv package for 'flexible' estimation:

```
1 > library(mgcv)
2 > gam1 <- gam(y~s(x), data = toy_data_02, subset = group=="Treated")
3 > gam0 <- gam(y~s(x), data = toy_data_02, subset = d==0); summary(gam0)
4
5 Family: gaussian
6 Link function: identity
7
8 Formula:
9 y ~ s(x)
10
11 Parametric coefficients:
12             Estimate Std. Error t value Pr(>|t|)
13 (Intercept)   0.2167      0.0307   7.059 2.05e-08 ***
14 ---
15 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
16
17 Approximate significance of smooth terms:
18             edf Ref.df    F p-value
19 s(x)         1      1 140.9 <2e-16 ***
20 ---
21 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
22
23 R-sq.(adj) = 0.782   Deviance explained = 78.8%
24 GCV = 0.03969   Scale est. = 0.037705   n = 40
```

Using GAMS




Using GAMS



Using GAMS

- We can estimate $\hat{\tau}_{\text{reg}}$ using the imputation estimator.

```
1 > cat("Estimate of ATE (GAM):",mean(predict(gam1) - predict(gam0)))  
2  
3 Estimate of ATE (GAM): 0.8379884
```



3/ Sensitivity Analysis

Cf. Cinelli & Hazlett (2020), *Making Sense of Sensitivity*, JRSS-B

Sensitivity Analysis for Regression

- Standard regression estimator of the ATE:

$$Y_i = \hat{\alpha} + \hat{\tau}D_i + \mathbf{X}'_i\hat{\beta} + \varepsilon_i$$

- Recall that regression with covariates both in uninteracted and interacted models are consistent for SATE/PATE (W4 discussion on Freedman 2008 critique and Lin 2023 remedy).

Sensitivity Analysis for Regression

- Standard regression estimator of the ATE:

$$Y_i = \hat{\alpha} + \hat{\tau}D_i + \mathbf{X}'_i\hat{\beta} + \varepsilon_i$$

- Recall that regression with covariates both in uninteracted and interacted models are consistent for SATE/PATE (W4 discussion on Freedman 2008 critique and Lin 2023 remedy).
- What if the true regression model contained U_i which we omitted?

$$Y_i = \alpha + \tau D_i + \mathbf{X}'_i\beta + \gamma U_i + \varepsilon_i$$

Sensitivity Analysis for Regression

- Standard regression estimator of the ATE:

$$Y_i = \hat{\alpha} + \hat{\tau}D_i + \mathbf{X}'_i\hat{\beta} + \varepsilon_i$$

- Recall that regression with covariates both in uninteracted and interacted models are consistent for SATE/PATE (W4 discussion on Freedman 2008 critique and Lin 2023 remedy).
- What if the true regression model contained U_i which we omitted?

$$Y_i = \alpha + \tau D_i + \mathbf{X}'_i\beta + \gamma U_i + \varepsilon_i$$

- Standard **omitted variable bias (OVB)** formula:

$$\begin{aligned}\hat{\tau} &= \tau + \gamma \times \underbrace{\frac{\text{cov}(D_i^{\perp \mathbf{X}}, U_i^{\perp \mathbf{X}})}{\mathbb{V}(D_i^{\perp \mathbf{X}})}}_{\text{regression of } U_i^{\perp \mathbf{X}} \text{ on } D_i^{\perp \mathbf{X}}} \\ &= \tau + \underbrace{\gamma \times \delta}_{\text{OVB}}\end{aligned}$$

- Gives us $\rightsquigarrow \tau = \hat{\tau} - \gamma\delta$

Partial R-Squared Interpretations

- Standard OVB written in terms of regression coefficients is difficult to reason about.
- Easier to reason with partial R^2 version of OVB (Cinelli and Hazlett, JRSSB 2020):

$$|\text{bias}| = \sqrt{\frac{R^2_{Y \sim U|D,\mathbf{x}} \cdot R^2_{D \sim U|\mathbf{x}}}{1 - R^2_{D \sim U|\mathbf{x}}} \cdot \frac{\mathbb{V}(Y \perp\!\!\!\perp \mathbf{x}, D)}{\mathbb{V}(D \perp\!\!\!\perp \mathbf{x})}}$$

Partial R-Squared Interpretations

- Standard OVB written in terms of regression coefficients is difficult to reason about.
- Easier to reason with partial R^2 version of OVB (Cinelli and Hazlett, JRSSB 2020):

$$|\text{bias}| = \sqrt{\frac{R^2_{Y \sim U|D, \mathbf{x}} \cdot R^2_{D \sim U|\mathbf{x}}}{1 - R^2_{D \sim U|\mathbf{x}}} \cdot \frac{\mathbb{V}(Y \perp\!\!\!\perp \mathbf{x}, D)}{\mathbb{V}(D \perp\!\!\!\perp \mathbf{x})}}$$

- Partial R^2 is the incremental predictive value of one variable:

$$R^2_{Y \sim U|D, \mathbf{x}} = \frac{R^2_{Y \sim D+\mathbf{x}+U} - R^2_{Y \sim D+\mathbf{x}}}{1 - R^2_{Y \sim D+\mathbf{x}}} = \frac{\text{additional variance explained by } U}{\text{variance unexplained by } D, \mathbf{x}}$$

Partial R-Squared Interpretations

- Standard OVB written in terms of regression coefficients is difficult to reason about.
- Easier to reason with partial R^2 version of OVB (Cinelli and Hazlett, JRSSB 2020):

$$|\text{bias}| = \sqrt{\frac{R^2_{Y \sim U|D, \mathbf{x}} \cdot R^2_{D \sim U|\mathbf{x}}}{1 - R^2_{D \sim U|\mathbf{x}}} \cdot \frac{\mathbb{V}(Y \perp\!\!\!\perp \mathbf{x}, D)}{\mathbb{V}(D \perp\!\!\!\perp \mathbf{x})}}$$

- Partial R^2 is the incremental predictive value of one variable:

$$R^2_{Y \sim U|D, \mathbf{x}} = \frac{R^2_{Y \sim D+\mathbf{x}+U} - R^2_{Y \sim D+\mathbf{x}}}{1 - R^2_{Y \sim D+\mathbf{x}}} = \frac{\text{additional variance explained by } U}{\text{variance unexplained by } D, \mathbf{x}}$$

- **Sensitivity analysis** can then vary two unknown parameters:
 - $R^2_{Y \sim U|D, \mathbf{x}} \in [0, 1]$: incremental predictive value of U for the outcome.
 - $R^2_{D \sim U|\mathbf{x}} \in [0, 1]$: incremental predictive value of U for treatment.
 - From these we can determine the bias and thus the true value of τ .


R Package: sensemakr

```
1 # load package
2 library(sensemakr)
3
4 # import dataset
5 data("darfur")
6
7 # run the regression model
8 model <- lm(peacefactor ~ directlyharmed + age + farmer_dar + herder_dar +
9           pastvoted + hhsize_darfur + female + village, data = darfur)
10
11 # conduct the sensitivity analysis
12 sensitivity <- sensemakr(model = model,
13                          treatment = "directlyharmed",
14                          benchmark_covariates = "female",
15                          kd = 1:3)
```

- female in our data is the strongest known confounder we are using as benchmark.
 - kd (knock-on distance): how many times stronger the confounder is related to the treatment in comparison to the observed benchmark covariate.

Sensitivity Analysis with Partial R^2

```
1 # plot bias contour of point estimate  
2 plot(sensitivity)
```



4/ Partial Identification

Cf. Manski (2003), *Partial Identification of Probability Distributions*

No-Assumption Bounds

- **Law of decreasing credibility** (Manski): credibility of inferences decreases with strength of assumptions
 - Idea: pick assumptions and then figure out what you can learn.
 - May not be point identified, but maybe we can **bound** the effect.

No-Assumption Bounds

- **Law of decreasing credibility** (Manski): credibility of inferences decreases with strength of assumptions
 - Idea: pick assumptions and then figure out what you can learn.
 - May not be point identified, but maybe we can **bound** the effect.
- If Y is bounded $[y_L, y_U]$, τ logically must be in $[y_L - y_U, y_U - y_L]$.

No-Assumption Bounds

- **Law of decreasing credibility** (Manski): credibility of inferences decreases with strength of assumptions
 - Idea: pick assumptions and then figure out what you can learn.
 - May not be point identified, but maybe we can **bound** the effect.
- If Y is bounded $[y_L, y_U]$, τ logically must be in $[y_L - y_U, y_U - y_L]$.
- Can we improve using data? Rewrite the ATE with $p = \mathbb{P}(D_i = 1)$:

$$\begin{aligned}\tau &= \mathbb{E}[Y_i \mid D_i = 1]p + \mathbb{E}[Y_i(1) \mid D_i = 0](1 - p) \\ &\quad - \mathbb{E}[Y_i(0) \mid D_i = 1]p - \mathbb{E}[Y_i \mid D_i = 0](1 - p)\end{aligned}$$

No-Assumption Bounds

- **Law of decreasing credibility** (Manski): credibility of inferences decreases with strength of assumptions
 - Idea: pick assumptions and then figure out what you can learn.
 - May not be point identified, but maybe we can **bound** the effect.
- If Y is bounded $[y_L, y_U]$, τ logically must be in $[y_L - y_U, y_U - y_L]$.
- Can we improve using data? Rewrite the ATE with $p = \mathbb{P}(D_i = 1)$:

$$\begin{aligned}\tau &= \mathbb{E}[Y_i \mid D_i = 1]p + \mathbb{E}[Y_i(1) \mid D_i = 0](1 - p) \\ &\quad - \mathbb{E}[Y_i(0) \mid D_i = 1]p - \mathbb{E}[Y_i \mid D_i = 0](1 - p)\end{aligned}$$

- Plug in y_L and y_U for the counterfactual means to get bounds for τ :

$$\tau \geq \mathbb{E}[Y_i \mid D_i = 1]p + y_L(1 - p) - y_U p - \mathbb{E}[Y_i \mid D_i = 0](1 - p)$$

$$\tau \leq \mathbb{E}[Y_i \mid D_i = 1]p + y_U(1 - p) - y_L p - \mathbb{E}[Y_i \mid D_i = 0](1 - p)$$

- These bounds have width of $|y_U - y_L|$ — half of the logical bounds.
- But always will contain 0. Weak assumptions $\rightsquigarrow$ weak inferences.

Optimized Treatment Choice

- **Assumptions** can narrow the bounds even further.

Optimized Treatment Choice

- **Assumptions** can narrow the bounds even further.
- Assumption: people choose the treatment with the highest outcome.
 - $\mathbb{E}[Y_i(0)|D_i = 1] \leq \mathbb{E}[Y_i(1)|D_i = 1] = \mathbb{E}[Y_i|D_i = 1]$
 - $\mathbb{E}[Y_i(1)|D_i = 0] \leq \mathbb{E}[Y_i(0)|D_i = 0] = \mathbb{E}[Y_i|D_i = 0]$

Optimized Treatment Choice

- **Assumptions** can narrow the bounds even further.
- Assumption: people choose the treatment with the highest outcome.
 - $\mathbb{E}[Y_i(0)|D_i = 1] \leq \mathbb{E}[Y_i(1)|D_i = 1] = \mathbb{E}[Y_i|D_i = 1]$
 - $\mathbb{E}[Y_i(1)|D_i = 0] \leq \mathbb{E}[Y_i(0)|D_i = 0] = \mathbb{E}[Y_i|D_i = 0]$
- Implies the following bounds for τ :

$$\tau \in [(1 - p)(y_L - \mathbb{E}[Y_i|D_i = 0]), p(\mathbb{E}[Y_i|D_i = 1] - y_L)]$$

- Width of these bounds: $\mathbb{E}[Y_i] - y_L$.
- Width now depends on the observed data!
- Interval will still always include zero.

Confidence Regions for Bounds

- More general setup:
 - True bounds $[\delta_L, \delta_U]$: the **identification region**.
 - Estimated bounds $[\hat{\delta}_L, \hat{\delta}_U]$.
 - $\widehat{\text{se}}(\hat{\delta}_L), \widehat{\text{se}}(\hat{\delta}_U)$: standard errors of the estimated bounds.

Confidence Regions for Bounds

- More general setup:
 - True bounds $[\delta_L, \delta_U]$: the **identification region**.
 - Estimated bounds $[\hat{\delta}_L, \hat{\delta}_U]$.
 - $\widehat{\text{se}}(\hat{\delta}_L), \widehat{\text{se}}(\hat{\delta}_U)$: standard errors of the estimated bounds.
- Two possible CI approaches that find intervals that...
 1. Covers the **identified region** with probability $1 - \alpha$:

$$\mathbb{P}(\hat{\delta}_L \leq \delta_L, \hat{\delta}_U \geq \delta_U) \geq 1 - \alpha$$

Confidence Regions for Bounds

- More general setup:
 - True bounds $[\delta_L, \delta_U]$: the **identification region**.
 - Estimated bounds $[\hat{\delta}_L, \hat{\delta}_U]$.
 - $\widehat{\text{se}}(\hat{\delta}_L), \widehat{\text{se}}(\hat{\delta}_U)$: standard errors of the estimated bounds.
- Two possible CI approaches that find intervals that...
 1. Covers the **identified region** with probability $1 - \alpha$:

$$\mathbb{P}(\hat{\delta}_L \leq \delta_L, \hat{\delta}_U \geq \delta_U) \geq 1 - \alpha$$

2. Covers the **true value of the parameter** with probability $1 - \alpha$:

$$\mathbb{P}(\tau \in [\hat{\delta}_L, \hat{\delta}_U]) \geq 1 - \alpha$$

Calculating Confidence Intervals

- Case 1: covering the identified region.

$$\left[\widehat{\delta}_L - z_{1-\alpha/2} \widehat{\text{se}}(\widehat{\delta}_L), \quad \widehat{\delta}_U + z_{1-\alpha/2} \widehat{\text{se}}(\widehat{\delta}_U) \right]$$

- Works because of **Bonferroni inequality**:

$$\mathbb{P}(\widehat{\delta}_L \leq \delta_L \text{ and } \widehat{\delta}_U \geq \delta_U) \geq \mathbb{P}(\widehat{\delta}_L \leq \delta_L) + \mathbb{P}(\widehat{\delta}_U \leq \delta_U) - 1 = 1 - \alpha$$

Calculating Confidence Intervals

- Case 1: covering the identified region.

$$\left[\widehat{\delta}_L - z_{1-\alpha/2} \widehat{\text{se}}(\widehat{\delta}_L), \quad \widehat{\delta}_U + z_{1-\alpha/2} \widehat{\text{se}}(\widehat{\delta}_U) \right]$$

- Works because of **Bonferroni inequality**:

$$\mathbb{P}(\widehat{\delta}_L \leq \delta_L \text{ and } \widehat{\delta}_U \geq \delta_U) \geq \mathbb{P}(\widehat{\delta}_L \leq \delta_L) + \mathbb{P}(\widehat{\delta}_U \leq \delta_U) - 1 = 1 - \alpha$$

- Case 2: covering the true parameter τ .

$$\left[\widehat{\delta}_L - z_{1-\alpha} \widehat{\text{se}}(\widehat{\delta}_L), \quad \widehat{\delta}_U + z_{1-\alpha} \widehat{\text{se}}(\widehat{\delta}_U) \right]$$

- If $\tau = \delta_L$ or $\tau = \delta_U$, coverage $\rightarrow 1 - \alpha$.
- If $\delta_L < \tau < \delta_U$, coverage $\rightarrow 1$.

Summing Up

- Conditional unconfoundedness is an assumption on unmeasured data, and thus inherently untestable!
 - The data are uninformative about the distribution of $Y(0)$ for treated units and $Y(1)$ for control units.
 - The assumption, however, can be 'indirectly' assessed. $\rightsquigarrow$ sensitivity analysis

Summing Up

- Conditional unconfoundedness is an assumption on unmeasured data, and thus inherently untestable!
 - The data are uninformative about the distribution of $Y(0)$ for treated units and $Y(1)$ for control units.
 - The assumption, however, can be 'indirectly' assessed. $\rightsquigarrow$ sensitivity analysis
- Without point identification, we can still learn from **bounds** under weaker assumptions.
 - Credibility–informativeness tradeoff (Manski).

Summing Up

- Conditional unconfoundedness is an assumption on unmeasured data, and thus inherently untestable!
 - The data are uninformative about the distribution of $Y(0)$ for treated units and $Y(1)$ for control units.
 - The assumption, however, can be 'indirectly' assessed. $\rightsquigarrow$ sensitivity analysis
- Without point identification, we can still learn from **bounds** under weaker assumptions.
 - Credibility–informativeness tradeoff (Manski).
- Up next: DAGs, i.e., a graphical framework for reasoning about which covariates to condition on.

Onto the presentations & discussions!

Contact Information:

jaewon.yoo@iss.nthu.edu.tw

<https://j1yoo4.github.io/>

