# PBA: 8. More on Regression

Fall 2026

Jaewon (“Jay-one”) Yoo

National Tsing Hua University

# Where are we? Where are we going?

- **Where we have been**
  - Part I, Working with Data in R (Weeks 1–3)
  - Part II, Causality, Prediction, and Regression (Weeks 5–10)
    - Week 5: causality. Causal effects, experiments, and observational studies
    - Week 6: relationships and tidying. Correlation, our own functions, joins
    - Week 7: prediction and iteration. Election polls, for loops, prediction errors and RMSE
    - Week 9: regression and model fit. Least squares, $R^2$, test error, bias and variance

# Where are we? Where are we going?

- **Where we have been**
  - Part I, Working with Data in R (Weeks 1–3)
  - Part II, Causality, Prediction, and Regression (Weeks 5–10)
    - Week 5: causality. Causal effects, experiments, and observational studies
    - Week 6: relationships and tidying. Correlation, our own functions, joins
    - Week 7: prediction and iteration. Election polls, for loops, prediction errors and RMSE
    - Week 9: regression and model fit. Least squares, $R^2$, test error, bias and variance
- **Today: more on regression, the last week of Part II**
  - Multiple regression; ridge, lasso, and elastic net; categorical predictors, an experiment, and interactions
  - After the lecture: E5 at your own pace with the TAs

# Where are we? Where are we going?

- **Where we have been**
  - Part I, Working with Data in R (Weeks 1–3)
  - Part II, Causality, Prediction, and Regression (Weeks 5–10)
    - Week 5: causality. Causal effects, experiments, and observational studies
    - Week 6: relationships and tidying. Correlation, our own functions, joins
    - Week 7: prediction and iteration. Election polls, for loops, prediction errors and RMSE
    - Week 9: regression and model fit. Least squares, $R^2$, test error, bias and variance
- **Today: more on regression, the last week of Part II**
  - Multiple regression; ridge, lasso, and elastic net; categorical predictors, an experiment, and interactions
  - After the lecture: E5 at your own pace with the TAs
- **Where we are going**
  - Part III, Statistical Uncertainty (Weeks 11–15). Next class (Nov 16): sampling and sampling distributions
  - Part IV, Showcase (Week 16)

# Roadmap

1. Multiple Regression
2. Many Predictors: Ridge, Lasso, and Elastic Net
3. Categorical Predictors and an Experiment
4. Interactions
5. In-Class Exercise E5

# 1/ Multiple Regression

# Beauty and Student Evaluations *Redux*

- Why do some instructors at universities receive high teaching evaluations scores from students while others receive lower ones?
    - Data contains instructor and course information on 463 courses taught at UT Austin (`ModernDive`; see recommended resources in syllabus).

| Variable | Description |
|---|---|
| `ID` | Identifier for courses taught at UT Austin |
| `score` | Avg. teaching score (1-5) |
| `bty_avg` | Instructor's avg. "beauty" score computed from a separate panel of six students (1-10). |
| `age` | Instructor's age |

Table: UT Austin Teaching Eval. Data

- ⇝ Q: Which better predicts student evaluation scores: an instructor's perceived attractiveness (beauty score) or their age?

# UT Austin Student Eval. Data

Import the data:

```
# Don't forget to load tidyverse and ggplot2!
> evals <- read_csv("https://bit.ly/46rSa7s"); evals

# A tibble: 463 × 14
      ID prof_ID score   age bty_avg gender ethnicity  language rank  pic_outfit pic_color cls_did_eval
   <dbl>   <dbl> <dbl> <dbl>   <dbl> <chr>  <chr>      <chr>    <chr> <chr>      <chr>             <dbl>
 1     1       1   4.7    36    5    female minority   english  tenu. not formal color              24
 2     2       1   4.1    36    5    female minority   english  tenu. not formal color              86
 3     3       1   3.9    36    5    female minority   english  tenu. not formal color              76
 4     4       1   4.8    36    5    female minority   english  tenu. not formal color              77
 5     5       2   4.6    59    3    male   not minor. english  tenu. not formal color              17
 6     6       2   4.3    59    3    male   not minor. english  tenu. not formal color              35
 7     7       2   2.8    59    3    male   not minor. english  tenu. not formal color              39
 8     8       3   4.1    51    3.33 male   not minor. english  tenu. not formal color              55
 9     9       3   3.4    51    3.33 male   not minor. english  tenu. not formal color             111
10    10       4   4.5    40    3.17 female not minor. english  tenu. not formal color              40
# ℹ 453 more rows
# ℹ 2 more variables: cls_students <dbl>, cls_level <chr>
# ℹ Use `print(n = ...)` to see more rows
```

# Multiple Predictors

What if we want to predict $Y$ as a function of many variables?

$$\texttt{score}_i = \alpha + \beta_1 \texttt{bty_avg}_i + \beta_2 \texttt{age}_i + \varepsilon_i$$

Why?

- Better predictions (at least in-sample).
- Better interpretation as **ceteris paribus** relationships:
  - $\beta_1$ is the relationship between `bty_avg` and `score` holding age constant.
  - **Statistical control** in a cross-sectional study.

Does this mean we should use all variables at our disposal?

- Share your thoughts on *eeclass*! (voluntary)

# Multiple Regression in *R*

```
> eval_cnt <- evals |> # Center the variables for better interpretation of the intercept
    mutate(bty_cnt = bty_avg - mean(bty_avg),
           age_cnt = age - mean(age))
> mult.fit.cnt <- lm(score ~ bty_cnt + age_cnt, data = eval_cnt); mult.fit.cnt

Call:
lm(formula = score ~ bty_cnt + age_cnt, data = eval_cnt)

Coefficients:
(Intercept)      bty_cnt      age_cnt
   4.174730     0.060656    -0.003059
```

- $\widehat{\alpha} \approx 4.17$: average score when both the beauty score and age are at their respective mean values (i.e., 4.42 & 48.4).
- $\widehat{\beta_1} \approx .061$: average increase in score for additional beauty score an instructor received, **holding age fixed**.
    - Ceteris paribus association b/w aesthetic evaluation and score, independent of age.
- $\widehat{\beta_2} \approx -.003$: average decrease in score for one-year increase in age, **holding beauty score fixed**.

Q: Centering the variable won't change the beta coefficients, but z-score transformation would. Why?

# Least Squares with Multiple Regression

- How do we estimate the coefficients?
- The same exact way as before: minimize prediction error!
- Residuals (aka prediction error) with multiple predictors:

$$Y_i - \widehat{Y}_i = \text{score}_i - \widehat{\alpha} - \widehat{\beta}_1 \text{bty_avg}_i - \widehat{\beta}_2 \text{age}_i$$

- Find the coefficients that minimize the sum of the squared residuals:

$$\text{SSR} = \sum_{i=1}^{n} (Y_i - \widehat{Y}_i)^2 = \sum_{i=1}^{n} (Y_i - \widehat{\alpha} - \widehat{\beta}_1 X_{i1} - \widehat{\beta}_2 X_{i2})^2$$

# Model Fit with Multiple Predictors

- $R^2$ mechanically increases when you add variables to the regression.
  - But this could be overfitting!
  - That is, our model may predict poorly on out-of-sample.
- Solution: penalize regression models with more variables.
  - Occam's razor: **simpler models are preferred**.
- Adjusted $R^2$: lowers regular $R^2$ for each additional predictor.

$$\text{Adjusted } R^2 = 1 - \underbrace{\frac{SSR}{TSS}}_{1-R^2} \times \frac{n-1}{n-k-1}$$

  - Week 9: $R^2 = \frac{TSS-SSR}{TSS} = 1 - \frac{SSR}{TSS}$. Here $n$: observations; $k$: predictors.
  - The factor $\frac{n-1}{n-k-1}$ grows with each predictor, so a new predictor must lower SSR enough.
  - If the added predictors don't help predict, adjusted $R^2$ goes down!

# Comparing Model Fits

- Model fit for the aesthetic evaluation model:

```
> fit.bty <- lm(score ~ bty_avg, data = evals)   # the beauty score model from Week 9
> library(broom)
> fit.bty |>
    glance() |>
    select(r.squared, adj.r.squared)

# A tibble: 1 × 2
  r.squared adj.r.squared
      <dbl>         <dbl>
1    0.0350        0.0329
```

- Model fit for multiple regression:

```
> mult.fit.cnt |>
    glance() |>
    select(r.squared, adj.r.squared)

# A tibble: 1 × 2
  r.squared adj.r.squared
      <dbl>         <dbl>
1    0.0378        0.0336
```

- $\rightsquigarrow$ Q: What if we observed the $adj.R^2$ decrease?

# Predicted Values from *R*

- Can we predict the evaluation score of a course currently being taught this semester (2026) with the following observed values?:

```
> mult.fit <- lm(score ~ bty_avg + age, data = evals)  # same model, not centered
> evals_26 <- evals |>   ## Suppose we have course 464 being taught in 2026
    select(ID, score, age, bty_avg) |>   # keep the four columns we use
    bind_rows(
      data.frame(ID = 464, score = NA, age = 36, bty_avg = 4.7)
    )
```

- The `modelr` package gives some functions that allow us to make predictions in a tidy way.
  - We could plug in values into the equation, but *R* can do this for us.
  - Let's use `add_predictions()` to predict the 2026 results:

```
> library(modelr)
> evals_26 |>
    filter(ID == 464) |>
    add_predictions(mult.fit)

# A tibble: 1 × 5
     ID score   age bty_avg  pred
  <dbl> <dbl> <dbl>   <dbl> <dbl>
1   464    NA    36     4.7  4.23
```

# Predictions from Several Models

- The `gather_predictions()` will return one row for each model passed to it with the prediction for that model:

```
> evals_26 |>
    filter(ID == 464) |>
    gather_predictions(fit.bty, mult.fit)

# A tibble: 2 × 6
  model       ID score   age bty_avg  pred
  <chr>    <dbl> <dbl> <dbl>   <dbl> <dbl>
1 fit.bty    464    NA    36     4.7  4.19
2 mult.fit   464    NA    36     4.7  4.23
```

# Predictions from New Data

- What about predicted values not in our data?
  - That is, what if we want to use our model to make predictions on other data?

```
> tibble(bty_avg = c(2.5, 4.5), age = 36) |>
    gather_predictions(fit.bty, mult.fit)

# A tibble: 4 × 4
  model    bty_avg   age  pred
  <chr>      <dbl> <dbl> <dbl>
1 fit.bty      2.5    36  4.05
2 fit.bty      4.5    36  4.18
3 mult.fit     2.5    36  4.10
4 mult.fit     4.5    36  4.22
```

- We can also get predicted values from the `augment()` function using the `newdata` argument:

```
> new_data <- tibble(bty_avg = c(2.5, 4.5), age = 36)
> augment(mult.fit, newdata = new_data)

# A tibble: 2 × 3
  bty_avg   age .fitted
    <dbl> <dbl>   <dbl>
1     2.5    36    4.10
2     4.5    36    4.22
```

# **2/** Many Predictors: Ridge, Lasso, and Elastic Net

# Predicting House Prices in Ames, Iowa

- Goal: predict the sale price of a house before it sells.
  - Data: 2,904 home sales in Ames, Iowa, from 2006 to 2010.[1]

| Variable | Description |
|---|---|
| `price` | Sale price in $1,000s (median: 160) |
| `Gr Liv Area` | Living area above ground (square feet) |
| `Overall Qual` | Overall quality of materials and finish (1–10) |
| `Year Built` | Year the house was built |
| `Garage Cars` | Size of the garage (number of cars) |
| … | 23 more numeric features (lot, baths, rooms, …) |

Table: Ames Housing Data

- ⇝ Q: Which features would you use to predict the price of a house?

[1] D. De Cock (2011), *Journal of Statistics Education* 19(3); Ames, Iowa Assessor's Office. Of the 2,930 sales in the file, 2,904 have all 27 numeric features.

# Ames Data in *R*: 40 Training Houses

- Import the data, keep the 27 numeric features, and draw 40 training houses:

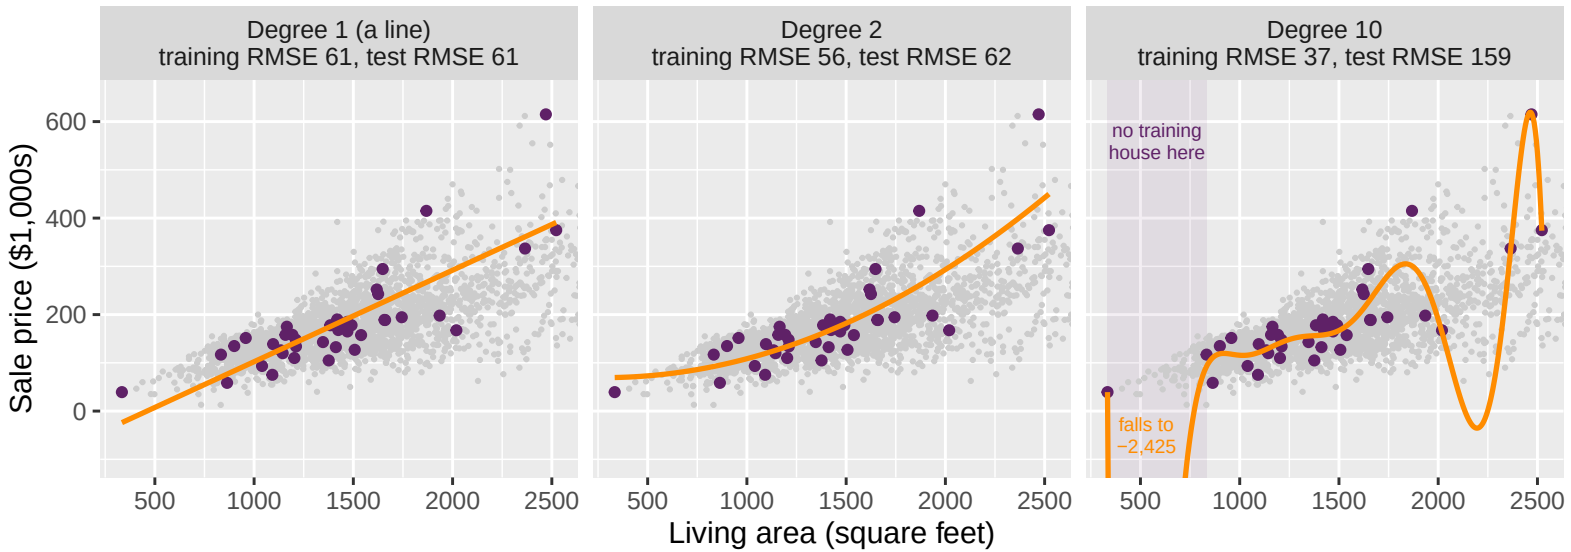
```
> ames <- read_tsv("https://jse.amstat.org/v19n3/decock/AmesHousing.txt") |>
    select(SalePrice, all_of(predictors)) |>   # predictors: 27 numeric features (R script)
    drop_na() |>                                # 2,904 sales with complete records
    mutate(price = SalePrice / 1000)            # sale price in $1,000s
> ames |> select(price, `Gr Liv Area`, `Overall Qual`, `Year Built`, `Garage Cars`) |>
    head(3)

# A tibble: 3 × 5
  price `Gr Liv Area` `Overall Qual` `Year Built` `Garage Cars`
  <dbl>         <dbl>          <dbl>        <dbl>         <dbl>
1   215          1656              6         1960             2
2   105           896              5         1961             1
3   172          1329              6         1958             1

> set.seed(2026)
> train <- sample(nrow(ames), 40)                          # 40 training houses
# data[rows, columns], as in Week 1 (R basics); -train means all rows but train
> ames_train <- ames[train, ]; ames_test <- ames[-train, ]  # test: the other 2,864 houses
```

- The next slides fit models to these 40 houses (few on purpose, slide 17) and test them on the other 2,864.

# Polynomial Regression Can Overfit




- Purple: the 40 training houses (slide 15). Grey: the test houses.[2]
- Degree 2: $\texttt{price}_i = \alpha + \beta_1 \texttt{area}_i + \beta_2 \texttt{area}_i^2 + \varepsilon_i$.
  - Linear regression adds up $\beta\times$ (a column). $\texttt{area}_i^2$ is one more column, so this polynomial is still linear regression: `lm(price ~ poly(area, 2))`.
- Degree 10 adds powers up to $\texttt{area}_i^{10}$:
  `lm(price ~ poly(` `` `Gr Liv Area` `` `, 10), data = ames_train)`.
  - It chases single houses: the best training fit and a much worse test fit (Week 9).

[2] Test RMSE over the 2,748 test houses inside the living-area range of the training houses. Polynomial regression: ISLR 2nd ed., Section 7.1.

# Few Houses: 27 Predictors Overfit

- In practice we train on most of the data (80/20). Here we use only 40 houses on purpose: are 27 predictors too many for so few houses?
- Each round: draw 40 houses at random, fit, and test on the other 2,864. The table shows the medians over 200 rounds:

```
> rmse <- function(y, yhat) sqrt(mean((y - yhat)^2))   # root mean squared error (Week 9)
> few <- c("Gr Liv Area", "Overall Qual", "Year Built")
# one round of the loop: the R script repeats the next four lines 200 times
> i <- sample(nrow(ames), 40)   # 40 random training houses
> fit_few <- lm(price ~ ., data = ames[i, c("price", few)])         # 3 predictors
> fit_all <- lm(price ~ ., data = ames[i, c("price", predictors)])  # all 27 predictors
> rmse(ames$price[-i], predict(fit_all, ames[-i, ]))   # test RMSE: the other 2,864 houses
```

| Least squares with | Training RMSE | Test RMSE |
|---|---|---|
| 3 predictors (living area, quality, year built) | 34.2 | 42.3 |
| All 27 predictors | 14.3 | 57.6 |

*Notes.* Price in $1,000s; medians over the 200 samples.

- With 27 predictors and 40 houses, least squares fits the noise.
  - The best training fit and the worst test fit (Week 9).

# Too Many Columns, Even with Many Houses

- Multiply every pair of the 27 predictors, for example living area × quality.
  - Like $\texttt{area}^2$ on slide 16, each product is one more column of numbers.
  - 27 predictors + 351 products = 378 columns.
- This time, train on many houses, as usual in practice: 2,323 (80%). Test on the other 581.
  - Repeat with 200 random splits. Code: today's R script.

| Least squares with | Training RMSE | Test RMSE |
|---|---|---|
| 27 predictors | 33.3 | 33.3 |
| 378 columns | 19.1 | 71.3 |

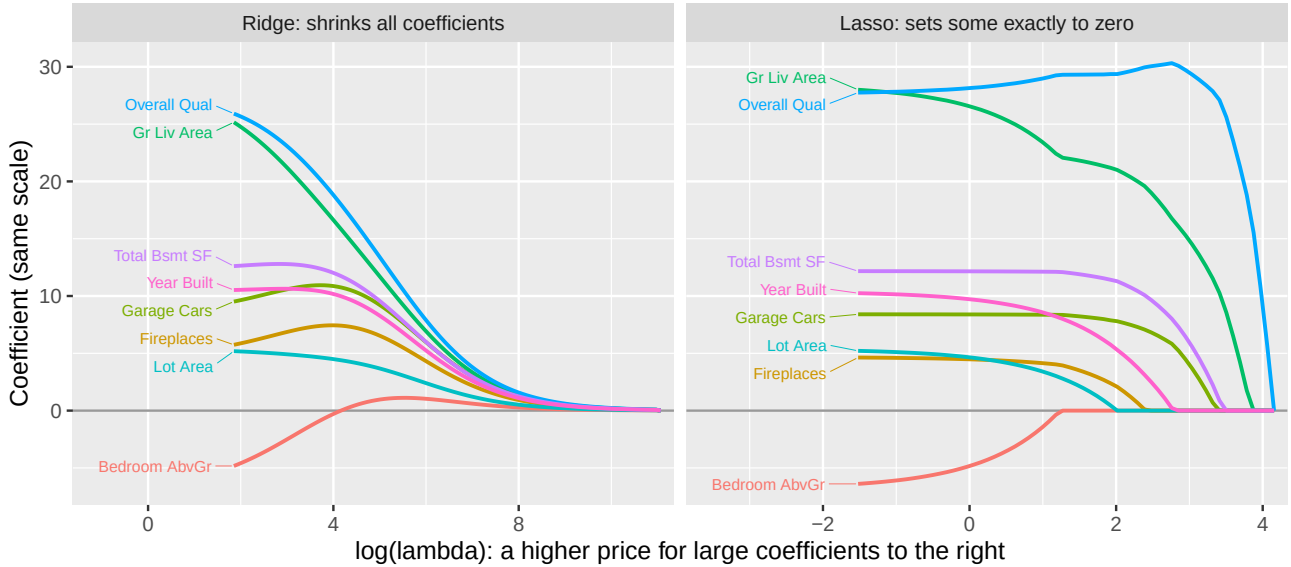
*Notes.* Price in $1,000s; medians over the 200 splits.

- Even with 2,323 training houses, the 378 columns fit the noise.
  - The training error falls, and the test error rises (in 172 of 200 splits).
  - In the worst split, the test RMSE was 12,133.
- With least squares, overfitting depends on how many coefficients we fit for the number of houses.
  - Too few houses (slide 17) or too many columns (here).

# Penalized Regression: A Price for Large Coefficients

- Least squares picks the coefficients with the smallest SSR.
  - With few houses per coefficient, some get very large to fit the noise.
- Penalized regression adds a price for large coefficients:
  - **Ridge**: SSR + $\lambda \times$ (sum of squared coefficients)
  - **Lasso**: SSR + $\lambda \times$ (sum of the absolute values of the coefficients)
  - **Elastic net**: a mix of the two prices.
- $\lambda$ (lambda) is the price.
  - $\lambda = 0$ gives least squares.
  - A very large $\lambda$ pushes every coefficient except the intercept to 0, so the model predicts the average.
  - A larger $\lambda$ gives a simpler model: a little more bias, much less variance (Week 9).
- The price is on the size of the coefficients.
  - Adjusted $R^2$ (slide 9) charges for the number of predictors instead.
  - The predictors are first put on the same scale, so the price does not depend on units (`glmnet` does this by default).

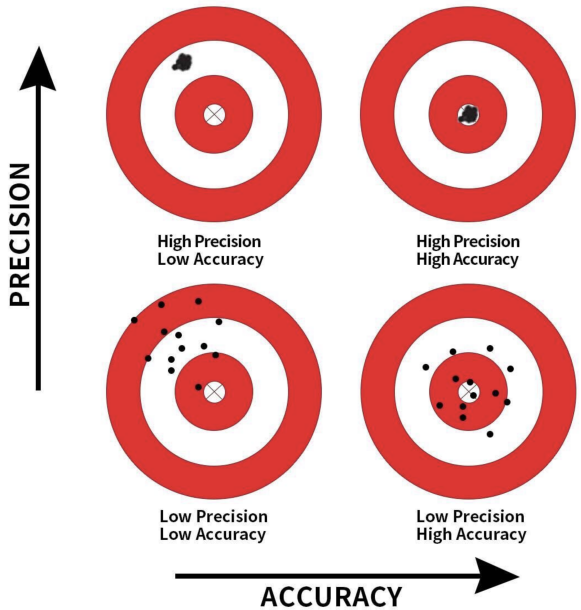
# Ridge Shrinks; Lasso Sets Some to Zero




- Moving right, $\lambda$ grows.
  - Ridge shrinks every coefficient but keeps them all: `glmnet(x, y, alpha = 0)`.
  - Lasso sets some exactly to 0, so it also chooses predictors: `alpha = 1`.
- Why exactly 0?[3]
  - Lasso charges the same price for each unit of a coefficient, so a weak predictor is not worth keeping at all.
  - Ridge charges very little for a small coefficient, so it keeps it.

[3] James, Witten, Hastie, and Tibshirani, *An Introduction to Statistical Learning*, 2nd ed., Sections 6.2.1–6.2.2.

# A Little Bias for Much Less Variance




Source: forecast.app

Next week we use the same target again. Each dot will then be an estimate from a different sample.

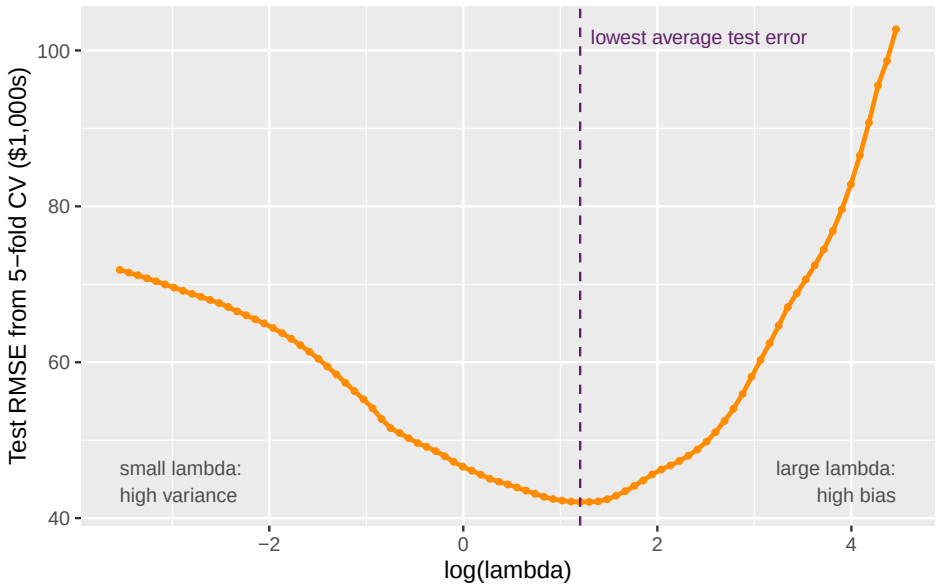
- Each dot: the predicted price of the same house from a model fit to a different set of 40 training houses (slide 17).
  - The center is the true value.
- Least squares with 27 predictors: the dots are far apart (high variance).
  - With 3 hand-picked predictors (slide 17), the dots stay close: a simple model has low variance.
- Ridge and lasso pull the dots together: much less variance.
  - They can add a little bias, because every coefficient is pulled toward 0.[4]

| Model | Spread ($1,000s) |
|---|---|
| Least squares, 3 predictors | 14 |
| Least squares, 27 predictors | 44 |
| Ridge | 21 |
| Lasso | 23 |
| Elastic net | 23 |

*Notes.* SD over 200 training sets, averaged over 900 test houses.

[4] ISLR 2nd ed., Section 6.2.1 and Figure 6.5. Code for the table: today's R script.

# Choosing $\lambda$: Cross-Validation




- Split the training data into 5 parts (folds).[5]
  - Fit on 4, test on the fifth, and rotate.
  - Average the 5 test errors for each $\lambda$, and pick the $\lambda$ with the lowest average.
- Each fold acts as a small test set, so we choose $\lambda$ without touching the 2,864 test houses.
  - Left: small $\lambda$, close to least squares (high variance). Right: large $\lambda$, close to the average (high bias).
- Here (one set of 40 training houses), the chosen lasso keeps 14 of the 27 predictors.

[5] ISLR 2nd ed., Section 5.1.3 (k-fold cross-validation). In *R*: `cv.glmnet()` (next slide). Code for the figure: today's R script.

# In *R*: Penalized Models Predict Better

- Choose $\lambda$ by 5-fold cross-validation (slide 22), then predict the test houses:

```
> library(glmnet)   # install.packages("glmnet") once
# ames_train, ames_test, predictors: made in today's R script
> x <- as.matrix(ames_train[, predictors])   # the 27 predictors
# alpha = 1: lasso; alpha = 0: ridge; between 0 and 1 (here 0.5): elastic net
> cv <- cv.glmnet(x, ames_train$price, alpha = 1, nfolds = 5)   # 5 folds (default: 10)
> coef(cv, s = "lambda.min")   # coefficients at the chosen lambda
> x_test <- as.matrix(ames_test[, predictors])
> predict(cv, newx = x_test, s = "lambda.min")
```

| 40 training houses, all 27 predictors | Median test RMSE |
|---|---|
| Least squares | 57.6 |
| Lasso | 42.3 |
| Ridge | 41.8 |
| Elastic net ($\alpha = 0.5$) | 41.2 |

- The penalized models do as well as our three hand-picked predictors (42.3), without our picking.
- Why elastic net? Garage cars and garage area carry almost the same information.
  - Lasso often kept only one of them (60% of 200 training sets).
  - Elastic net kept both more often (44% vs. 26%).[6]

[6] Training sets of slide 21. Same point in An Introduction to glmnet.

# Three Test Houses: Who Predicts Closer?

| | House A | House B | House C |
|---|---|---|---|
| Living area (square feet) | 896 | 1,145 | 1,836 |
| Overall quality (1–10) | 5 | 7 | 8 |
| Year built | 1961 | 2006 | 2004 |
| **Actual price** | **105** | **160** | **280** |
| Least squares, 3 predictors | 107 | 214 | 301 |
| Least squares, 27 predictors | 291 | 184 | 268 |
| Lasso | 191 | 194 | 258 |
| Ridge | 170 | 185 | 273 |
| Elastic net | 189 | 195 | 262 |

*Notes.* Three test houses: the 10th, 50th, and 90th percentiles of price. Prices in $1,000s.

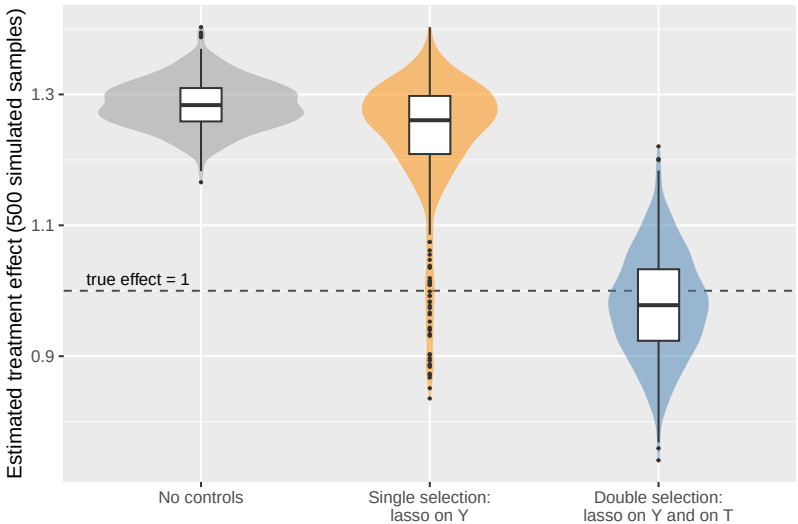
Every model is fit to the same 40 training houses (slide 15). Code: today's R script.

- Prediction gives a price for each new house. The error is the actual price minus the predicted price.
- House A sold for 105.
    - Least squares with 27 predictors predicts 291 (high variance).
    - The penalized models predict 170 to 191: closer, but still far.
- No model is right for every house.
    - Test RMSE (slide 23) summarizes the errors over all 2,864 test houses.

# Careful: Lasso Picks Predictors, Not Causes

- Lasso keeps the variables that help predict the outcome, whatever the reason they help.
    - So it cannot tell a **confounder**, which we must control for, from a variable that the treatment or the outcome itself changes, which we must not control for.
- A zero is not “no effect” either.
    - It means “not needed once the other predictors are in”: lasso often dropped garage cars when it kept garage area (slide 23).
- To learn what a variable does, we need a fair comparison, such as an experiment (later today).

# Lasso for Controls: Use Double Selection




- Simulated: 200 people, 50 candidate controls. True effect: 1.
  - 5 confounders drive $T$ strongly and $Y$ weakly.
- **Single selection** (lasso on $Y$, $T$ kept in): estimate 1.23.
  - Given $T$, the confounders barely help predict $Y$, so lasso drops them (keeps about 1 of 5).
- **Double selection**: also run lasso of $T$ on the controls, and control for every variable chosen in either step: estimate 0.98.
  - Keeps all 5.[7]
- To choose controls, ask two questions: what predicts $Y$, and what predicts $T$?

[7] Belloni, Chernozhukov, and Hansen (2014), *Review of Economic Studies* 81(2): 608–650. Simulation: 500 samples; $\lambda$ by 5-fold CV (`lambda.min`). Code: today's R script.

# 3/ Categorical Predictors and an Experiment

# Binary Predictors

$$Y_i = \alpha + \beta X_i + \varepsilon_i$$

- When the predictor $X_i$ is **binary**:
  - Intercept $\widehat{\alpha}$ is the average outcome in the $X = 0$ group.
  - Slope $\widehat{\beta}$ is the difference-in-means of $Y$ between $X = 1$ group and $X = 0$ group.

$$\widehat{\beta} = \overline{Y}_{\text{treated}} - \overline{Y}_{\text{control}}$$

- If there are other predictors, this becomes the difference in means controlling for them.

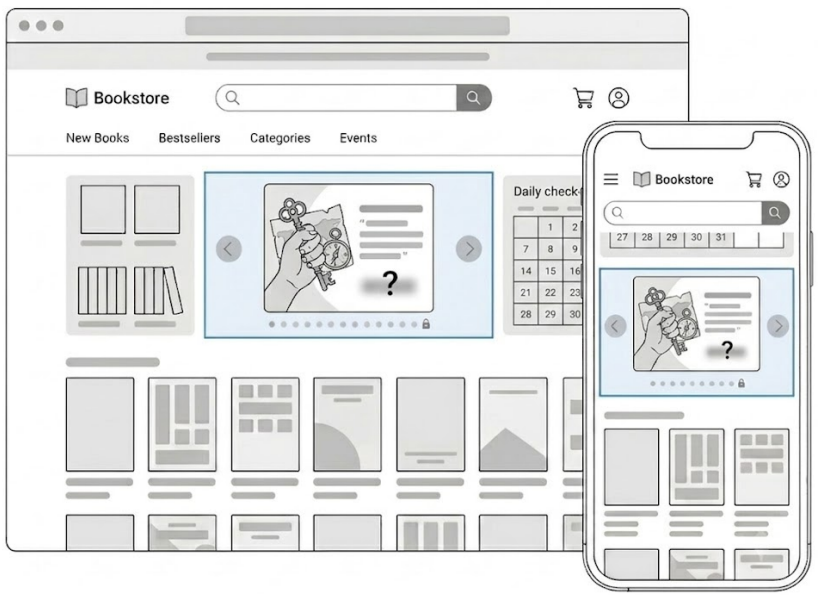
# A Research Example: Blind Lead-In (BLI) Ads




- An interactive ad at a large Korean online bookstore: swipe through 6 to 29 panels (median 12).
  - The book is shown only on the last panel. The paper calls this a blind lead-in (BLI) ad.
- Viewers swipe at their own pace and can stop at any panel.
  - Only those who reach the end see the book.[8]

[8] Yoon, Yoo, Kim, Kim, and Lee, "Interactive Blind Lead-In Ads and Purchase Conversion," *Journal of Interactive Marketing*, forthcoming.

# A Field Experiment: Slider Ad vs. Banner




Schematic of the bookstore's daily check-in page, where the ads ran.

- Same study: 109,156 customers, 80 days, 76 new books.
- Each customer was randomly assigned to see:
  - the **slider ad** (book shown at the end), or
  - a **static banner** (book shown at once).
- Outcome: did the customer buy the advertised book within 30 days? (1 or 0)
- Q: Does the slider ad sell more books than the banner?

# Ad Experiment Data (Synthetic)

| Variable Name | Description |
|---|---|
| `treat_ind` | 1 if assigned to the slider ad, 0 if assigned to the banner |
| `finished` | 1 if the customer swiped to the last panel (always 0 for the banner) |
| `purchase` | 1 if the customer bought the advertised book within 30 days |

```
# today's data are in the course data package pbadata (update it once: see the top of the R script)
> library(pbadata)
> bli <- bli_experiment   # one row per customer visit: 4,000 synthetic customers, 76 days
> bli |> head()
# A tibble: 6 × 6
  customer date         book treat_ind finished purchase
     <int> <date>      <int>     <int>    <int>    <int>
1        8 2017-06-13 100001         0        0        0
2        9 2017-06-13 100001         0        0        0
3       18 2017-06-13 100001         1        0        0
4       20 2017-06-13 100001         1        0        0
5       27 2017-06-13 100001         0        0        0
6       29 2017-06-13 100001         0        0        0
```

Data: a synthetic version of the experiment, from the study's replication package (deposited in Harvard Dataverse; public release forthcoming). The real customer records are under a nondisclosure agreement (NDA) with the retailer, so the file is simulated to keep the directions of the real results. The numbers differ from the paper, and purchase rates are much higher than in the real data.

# Linear Regression for Experimental Data

- Difference in means by assignment:

```
> bli |>
    group_by(treat_ind) |>
    summarize(purchase = mean(purchase)) |>
    pivot_wider(names_from = treat_ind, values_from = purchase) |>
    mutate(ITT = `1` - `0`)
# A tibble: 1 × 3
     `0`    `1`     ITT
   <dbl>  <dbl>   <dbl>
1 0.0127 0.0143 0.00164
```

- The same number from `lm()`:

```
> lm(purchase ~ treat_ind, data = bli) |> coef() |> round(5)
(Intercept)   treat_ind
    0.01266     0.00164
```

- Same `lm()` as for prediction.
  - Under random assignment, the coefficient on `treat_ind` is the effect of being assigned the slider ad, the **intent-to-treat (ITT)** effect.
  - +0.16 percentage points (1.27% to 1.43%).

# Who Finishes the Ad?

- Among visits assigned to the slider ad, compare those that finished with those that did not:

```
> bli |>
    filter(treat_ind == 1) |>
    group_by(finished) |>
    summarize(n = n(), purchase = mean(purchase))
# A tibble: 2 × 3
  finished     n purchase
     <int> <int>    <dbl>
1        0 36750  0.00955
2        1  1242  0.155
```

- Only 1,242 of 37,992 slider visits (3.3%) reached the last panel.
- Visits that finished bought far more: 15.5% vs. 1.0%.
    - A gap of 14.5 percentage points.
- Q: Is 14.5 points the effect of finishing the ad?

# Compare by Assignment, Not by Choice

- Viewers chose to finish.
  - Those who finish may be likely buyers anyway, so comparing groups that people choose into is not fair (Week 5).
- Random assignment gives the fair comparison. Few viewers finish, so the effect of assignment is diluted:

$$\text{effect of assignment} \approx \text{effect of finishing} \times \text{share who finish}$$

  - This needs one assumption: for viewers who would quit the slider ad early, the slider ad and the banner give the same chance of buying.
- So the effect of finishing $\approx 0.00164/0.0327 \approx 0.050$: about 5 percentage points.
  - A third of the 14.5-point gap.
  - It is the effect for viewers who finish, compared with the banner.
- The real data show the same pattern: finishing has a much larger effect than assignment, because few people finish.

# How Can We Explain Finishing?

- Finishing matters: the book is shown only at the end, and finishing has a much larger effect than assignment (slide 33).
- Yet few viewers finish: 3.3% of all views, from 0.7% to 7.7% across the 74 ads. Why do some ads keep viewers going?
- The paper starts from theory. At each panel, a viewer asks: is the next panel worth the effort?
  - **Value**: more text on the first panel tells viewers what the story is about (information scent). So more viewers finish.
  - **Effort**: more text and more colorful images on the late panels take effort. So fewer viewers finish.
  - **Faces**: a human face before the last panel adds social and emotional value. So more viewers finish.
- Here the goal is to explain, not to predict: theory picks the variables, and the data check the signs.

# Ad Data: 74 Slider Ads

| Variable Name | Description |
|---|---|
| `completion` | % of views that reached the last panel |
| `first_chars` | Number of characters of text on the first panel |
| `late_chars` | Characters of text per panel on the late panels (before the last) |
| `late_color` | How colorful the images on the late panels are |
| `middle_face` | 1 if a human face appears on the middle panels |
| … | 14 more: date, number of panels, tone and story of the text, faces, author photo |

```
> bli_ads |> select(ad, completion, first_chars, late_chars, late_color, middle_face) |> head()
# A tibble: 6 × 6
  ad    completion first_chars late_chars late_color middle_face
  <chr>      <dbl>       <int>      <dbl>      <dbl>       <int>
1 Ad_01      0.679          23       35.6      0.360           1
2 Ad_02      1.77           21       39.1      0.453           1
3 Ad_03      3.00           21       32.6      0.524           1
4 Ad_04      4.74           79       94.8      0.433           1
5 Ad_05      1.48           47       75.8      0.592           1
6 Ad_06      1.88           41       45.8      0.537           0
```

Data: one row per ad, from the study's replication package (real data; book titles removed).

# Regression on 74 Ads: Do the Signs Match?

```
> ads <- bli_ads |>  # completion: % of views that reached the last panel
    # z-scores of the 13 continuous features (1 = one standard deviation); middle_face stays 1 or 0
    mutate(across(total_slides:late_story, ~ (.x - mean(.x)) / sd(.x)))
> lm(completion ~ first_chars + late_chars + late_color + middle_face, data = ads) |>
    coef() |> round(3)
(Intercept) first_chars  late_chars  late_color middle_face
      3.363       0.189      -0.158      -0.358       0.396
```

- The signs follow the theory, each holding the other features fixed (ceteris paribus, slide 6):
  - More text on the first panel: higher completion (scent).
  - More text and more colorful images on the late panels: lower completion (effort).
  - A face on the middle panels: higher completion (by 0.4 percentage points).
- Ad features were not randomly assigned, so these are associations, not causal effects (Week 5).

# The Survey Data

- A follow-up survey for the same study: 240 participants each viewed 10 slider ads. They could quit an ad at any panel.
- For each ad:
  - `finished`: 1 if they swiped to the last panel (recorded, not asked)
  - `experience`: asked after the ad, agreement with "I am having a good time right now," 1 (strongly disagree) to 5 (strongly agree)
  - `purchase_intent`: how likely are they to buy the book? 1 to 5
  - `genre_match`: 1 if the book's genre is one they prefer
- Data: de-identified responses from the same replication package. Rows for books a participant had already read are dropped.

```
> survey <- bli_survey   # one row per participant and ad
> survey |> head()
# A tibble: 6 × 6
  person  book finished experience purchase_intent genre_match
   <int> <int>    <int>      <int>           <dbl>       <int>
1      1     2        1          3               2           0
2      1     3        1          4               3           0
3      1     6        1          3               2           0
4      1    10        1          1               1           0
5      1    11        1          2               2           0
6      1    17        0          1               1           0
```

# Categorical Variables in Regression

- We often have **categorical variables**, e.g.,:
  - Experience of viewing an ad: 1 (strongly negative), 2, 3, 4, 5 (strongly positive)
  - Book advertised: 76 titles
- How to incorporate into regression: create a **series of binary variables** (also called **dummy variables**)

| Unit | Experience | Exp. 1 | Exp. 2 | Exp. 3 | Exp. 4 | Exp. 5 |
|---|---|---|---|---|---|---|
| 1 | 1 | 1 | 0 | 0 | 0 | 0 |
| 2 | 2 | 0 | 1 | 0 | 0 | 0 |
| 3 | 3 | 0 | 0 | 1 | 0 | 0 |
| 4 | 4 | 0 | 0 | 0 | 1 | 0 |
| 5 | 5 | 0 | 0 | 0 | 0 | 1 |
| ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ | ⋮ |

- Then include **all but one** (i.e., baseline) of these binary variables:

$$\text{intent}_i = \alpha + \beta_1 \text{Exp2}_i + \beta_2 \text{Exp3}_i + \beta_3 \text{Exp4}_i + \beta_4 \text{Exp5}_i + \varepsilon_i$$

# Interpreting Categorical Variables

$$\text{intent}_i = \alpha + \beta_1 \text{Exp2}_i + \beta_2 \text{Exp3}_i + \beta_3 \text{Exp4}_i + \beta_4 \text{Exp5}_i + \varepsilon_i$$

- $\widehat{\alpha}$: average purchase intention in the omitted group/baseline (experience 1, strongly negative).
- $\widehat{\beta}$ coefficients: average difference between each group and the baseline.
  - $\widehat{\beta_1}$: average difference in purchase intention between views rated 2 and views rated 1.
  - $\widehat{\beta_2}$: average difference in purchase intention between views rated 3 and views rated 1.

- Q: how do we interpret $\widehat{\beta_4}$?

# Regression with Categorical Variables

- Convert the categorical variable into a `factor`; the first level is the baseline. Without `factor()`, *R* treats 1 to 5 as numbers and fits one straight line:

```
> survey <- survey |>   # levels 1 to 5; level 1 (strongly negative) is the baseline
    mutate(experience = factor(experience))
```

- Regress `purchase_intent` on the factor variable:

```
> fit.exp <- lm(purchase_intent ~ experience, data = survey); fit.exp

Call:
lm(formula = purchase_intent ~ experience, data = survey)

Coefficients:
(Intercept)  experience2  experience3  experience4  experience5
      1.170        0.658        1.461        2.109        2.911
```

- A better experience goes with higher purchase intention: views rated 5 score 2.9 points higher than views rated 1 (on a 1–5 scale).

# What *R* Does Internally w/ Factors in `lm()`

- The factor/categorical variable, `experience`:

```
> survey |> select(experience) |> pull() |> head()
[1] 3 4 3 1 2 1
Levels: 1 2 3 4 5
```

- Notice, *R* automatically created a set of indicator variables for each level except the baseline:

```
> model.matrix(fit.exp) |> head()
  (Intercept) experience2 experience3 experience4 experience5
1           1           0           1           0           0
2           1           0           0           1           0
3           1           0           1           0           0
4           1           0           0           0           0
5           1           1           0           0           0
6           1           0           0           0           0
```

# Back to the Experiment: Book Dummies as Controls

- Add the book advertised on each day (76 books) as a categorical predictor:

```
> lm(purchase ~ treat_ind + factor(book), data = bli) |> coef() |> head(2) |> round(5)
(Intercept)   treat_ind
    0.01714     0.00165
```

- `factor(book)` adds 75 binary variables (76 books, one baseline).
- Under randomization, the control barely moves the estimate: 0.00164 to 0.00165.
    - Randomization, not the control, makes the comparison fair.

# 4/ Interactions

# Interactions: When One Effect Depends on Another

- Does the effect of one predictor depend on another?
  - Is the gap in purchase intention between finished and unfinished ads larger when viewing felt good?
- Add the product of the two predictors, an **interaction**:

$$Y_i = \alpha + \beta_1 X_i + \beta_2 Z_i + \beta_3 (X_i \times Z_i) + \varepsilon_i$$

  - The effect of $X$ is $\beta_1$ when $Z = 0$, and $\beta_1 + \beta_3$ when $Z = 1$.
  - With a categorical $Z$, there is one interaction for each level except the baseline. In *R*: `y ~ x * z`.

# Finishing × Experience in *R*

- The paper's model (Table 8): `finished` and its interaction with each experience level. Person and book dummies compare views of the same person and of the same book:

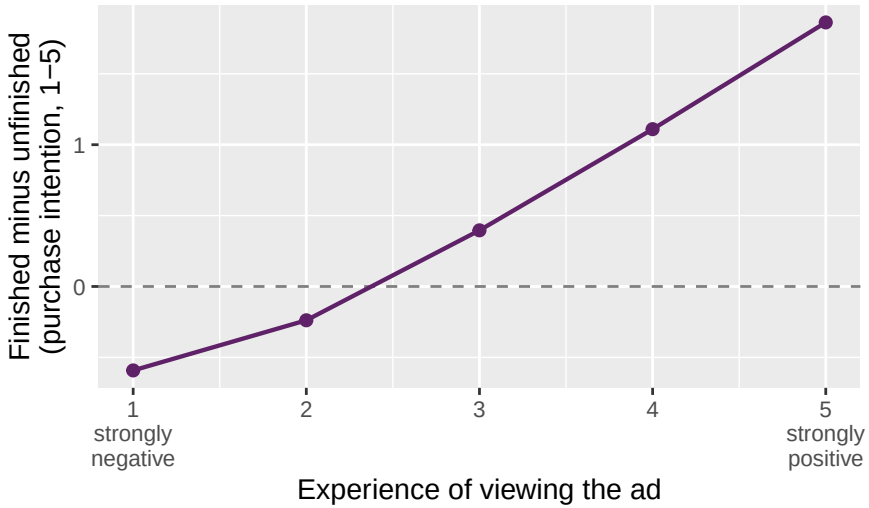
```
> fit.int <- lm(purchase_intent ~ finished + finished:experience + genre_match +
                  factor(person) + factor(book), data = survey)
> coef(fit.int)[c("finished", paste0("finished:experience", 2:5))] |> round(3)
            finished finished:experience2 finished:experience3
              -0.592                0.354                0.988
finished:experience4 finished:experience5
               1.702                2.454
```

- Gap between finished and unfinished ads at experience level $k$ = `finished` + `finished:experience`$k$:

| Experience | 1 | 2 | 3 | 4 | 5 |
|---|---|---|---|---|---|
| Finished − unfinished | −0.59 | −0.24 | +0.40 | +1.11 | +1.86 |

- Finishing goes with higher purchase intention only when viewing the ad felt good.

# Plotting the Interaction




- Each point is the finished-minus-unfinished gap at one experience level (slide 44):

```
> effects <- function(fit) {   # the effect of finishing at each experience level
    b <- coef(fit)
    extra <- b[paste0("finished:experience", 2:5)]   # extra effect at levels 2 to 5
    tibble(experience = 1:5, effect = b[["finished"]] + c(0, extra))
  }
> ggplot(effects(fit.int), aes(x = experience, y = effect)) + geom_line() + geom_point()
```

- Without the interaction, the model would force the same gap at every level: a flat line.

# 5/ In-Class Exercise E5

# In-Class Exercises: Background

- You are working for an investor who is thinking of building a new condominium in NYC and wants to predict the market value before making any decision.
    - The new condominium will be located in Manhattan and will be 220,000 square feet large with 300 units.

- The NYC `housing` dataset contains some information about the condominiums located in New York City and thus can be used to make the prediction. Some relevant variables:
    - `Market.Value.per.SqFt`: total market value per square foot of the land and building
    - `Boro`: the location (borough) of the condominium in NYC (e.g., Manhattan)
    - `Total.Units`: total number of units in the building
    - `Gross.SqFt`: the size of the building in square foot

# In-Class Exercises: Task

- Predict the market value per square foot of the new condominium with two regression models:
  - Model 1: the location (`boro`) only.
  - Model 2: the location, the size (`SqFt`), and the number of units (`units`).
- Compare the two models, and tell the investor whether to build.
  - To get started, we'll give you a few clues on the following slide.

# In-Class Exercises: Instructions

1. Import the data using `read_csv()` from `"https://bit.ly/49QWVuc"` and store the resulting tibble as `housing`. Using `select()` and `rename()`, keep 'Market.Value.per.SqFt', 'Boro', 'Total.Units', and 'Gross.SqFt' renamed as 'value', 'boro', 'units', and 'SqFt', respectively, for the subsequent analysis. Don't forget to overwrite `housing` with the subsetted data.
2. Specify two regression models where you consider 1) only the location (i.e., `boro`) as the predictors and 2) the location as well as the size (i.e., `SqFt`) & units (i.e., `units`) as the predictors. Write out the equations for the two models.
3. Estimate the models using `lm()` and interpret the coefficients.
4. Compare the 'fit' of the two models by looking at the adjusted $R^2$ with `glance()`. That is, do size and number of units 'add value' to our prediction?
5. Predict the market value of the condominium using `modelr` functions. What is the value of the condominium predicted to be given the specifications on slide 46?
6. The investor wants to build only if the valuation is predicted to be higher than 140 per Sq Ft. Should you recommend investing or not?

# What Comes Next?

- Today: multiple regression; ridge, lasso, and elastic net; categorical predictors; regression for an experiment; interactions; difference-in-differences as a regression.
- Next class (Nov 16): **sampling and sampling distributions**. Complete the Required reading in the weekly schedule beforehand.
- Due Sunday, Nov 15: Problem Set 1 and Milestone 3 (first visualization). Assignment 3 is due before class on Monday, Nov 16. Problem Set 2 (IMS Ch. 8) is released today.
- E5 (today's exercise): submit on eeclass by 23:59 today.

# Enjoy Your Weekends! :)

*Contact Information:*
jaewon.yoo@iss.nthu.edu.tw
https://j1yoo.github.io/



# Difference-in-Differences as a Regression

- In Week 7, we computed DiD by hand with the minimum wage data: about 2.45 workers per restaurant.
- An interaction (slide 43) gives the same number:

$$Y_{it} = \alpha + \beta_1 \text{NJ}_i + \beta_2 \text{Post}_t + \beta_3 (\text{NJ}_i \times \text{Post}_t) + \varepsilon_{it}$$

```
> minwage <- as_tibble(read.csv(url("https://bit.ly/3QZmodH")))   # same data as before
> long <- minwage |>
    mutate(id = 1:n(), nj = if_else(location == "PA", 0, 1),
           fte_before = fullBefore + 0.5 * partBefore,
           fte_after = fullAfter + 0.5 * partAfter) |>
    pivot_longer(c(fte_before, fte_after), names_to = "period", values_to = "fte") |>
    mutate(post = if_else(period == "fte_after", 1, 0))
> lm(fte ~ nj * post, data = long) |> coef() |> round(2)
(Intercept)          nj        post     nj:post
      20.30       -3.01       -2.13        2.45
```

| | |
|---|---|
| `(Intercept)` | Average FTE in PA before the law |
| `nj` | NJ minus PA before the law |
| `post` | PA's change: the change we predict for NJ without the law |
| `nj:post` | NJ's change minus PA's change: the DiD |

# Two-Way Fixed Effects

- **Fixed effects** are dummies. **Two-way fixed effects (TWFE)** use one set of dummies for the units and one set for the time periods:

$$Y_{it} = \alpha_i + \lambda_t + \beta_3(\text{NJ}_i \times \text{Post}_t) + \varepsilon_{it}$$

- $\alpha_i$: one intercept for each restaurant. $\lambda_t$: one intercept for each period. Same idea as `factor(person) + factor(book)` on slide 44.

```
> fit.twfe <- lm(fte ~ nj:post + post + factor(id), data = long)
> coef(fit.twfe)[c("post", "nj:post")] |> round(2)
   post nj:post
  -2.13    2.45
```

- `factor(id)` adds 357 binary variables (358 restaurants, one baseline).
  - We leave out `nj` because the restaurant dummies already contain it.
- DiD is the design. TWFE is one estimator.
  - It still needs parallel trends.
  - All NJ restaurants are treated at one date. With different dates, TWFE is a weighted average of many DiD comparisons, and some use already-treated units as controls (Goodman-Bacon, 2021).